# How Many Demonstrations Does a €400 Robot Arm Need?

Low training loss, low task success: diagnosing a perceptual bottleneck in imitation learning on the SO-101 arm

Fabio Bariletti

MSc Artificial Intelligence Systems · July – October 2026

Code · Datasets & models · Videos

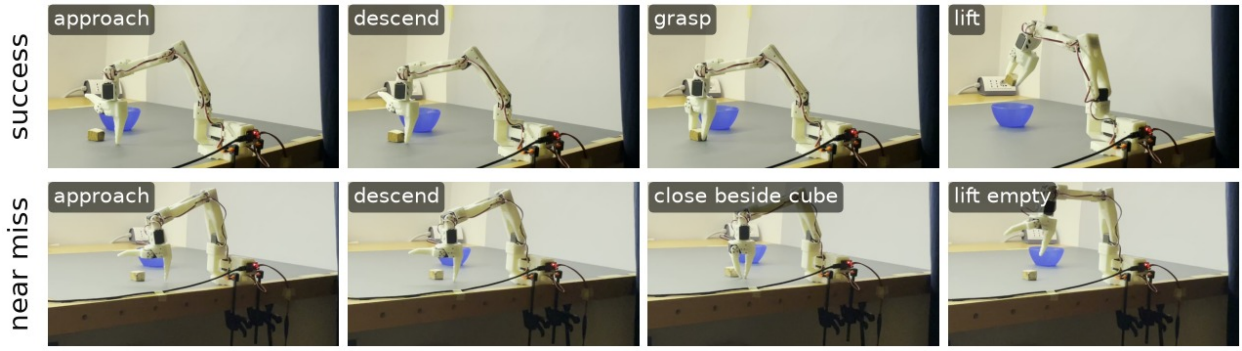



**Figure 1:** Autonomous rollouts of ACT policies trained on setup-2 data. **Top:** success (`v2_100`). **Bottom:** the characteristic failure, a *near miss* (`v2_024`) — the policy executes a well-formed grasp in the right place, but closes within about a centimetre of the cube and lifts an empty gripper.

**4× data**
24 → 100 demonstrations: success 1/15 → 2/15, **no detectable gain**

**≈1 cm**
8–12 of 15 rollouts per model are **near misses**: right motion, gripper closes next to the cube

**Likely cause**
at the grasp, neither camera shows gripper and cube together: the limit appears **perceptual** rather than data-related

**Abstract**

I built an SO-101 robot arm from parts, recorded three datasets of teleoperated demonstrations across two camera setups, trained four ACT policies, and evaluated three of them on a pick-and-place task under a protocol fixed in advance. **The headline result is negative, and the interesting part is why.** All policies converge in training (final $L_1$ loss 0.05–0.09, an order of magnitude below its initial value) but succeed on only 7–13% of attempts. Quadrupling the data from 24 to 100 demonstrations produces no measurable improvement, and a redesign of cameras and visual contrast does not raise the success rate either. Yet the policies are far from random: most failures are near misses within about a centimetre of the cube, and in longer rollouts the arm notices a missed grasp and goes back to retry. Training loss, outcome statistics and the recorded observations point to a **perceptual bottleneck**: neither camera captures gripper and object together at the moment of grasping, so the policy knows roughly where the cube is but cannot resolve the final alignment. An earlier data-level failure — a gripper that never opened — was traced to the action distribution of the demonstrations and fixed by changing how they were recorded. All datasets and checkpoints are public.

## 1 Introduction

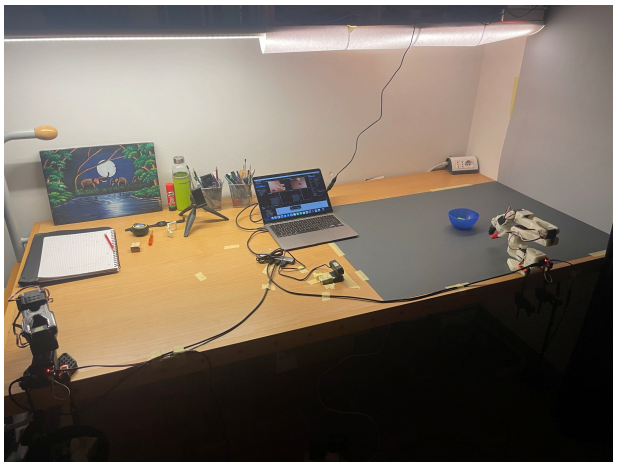
Behaviour cloning on low-cost hardware has become accessible: the SO-101 arm costs around €400 in parts, and Hugging Face's LeRobot provides a complete pipeline from teleoperated data collection to policy deployment. Many hobbyist projects now show a trained arm picking up a cube. Far fewer report *measured* success rates, confidence intervals, or an analysis of *why* a policy fails. This project treats the arm as an experimental platform rather than a demo, and asks:

> *How many demonstrations does an ACT policy need to learn a manipulation task on low-cost hardware — and what limits performance when it does not?*

**Contributions.**

- A controlled data-efficiency comparison (24 vs. 100 demonstrations, nested subsets) at a fixed training budget, reported with Wilson confidence intervals and a graded, per-position outcome scale.
- A diagnosis of the performance ceiling as a perceptual bottleneck, supported by training loss, outcome statistics, recorded observations and an intervention on the camera setup.
- Evidence that the policies nonetheless learned the task structure, including closed-loop recovery from a missed grasp.
- A data-distribution failure (gripper action imbalance) identified from dataset statistics and fixed by changing demonstration technique rather than by training longer.
- Public datasets, model checkpoints and figure-generation code.

## 2 Platform and task



| | |
|---|---|
| Robot | SO-101 leader–follower pair |
| Actuators | 12× Feetech STS3215 servos |
| | follower: 6× 12 V, 1:345 |
| | leader: 7.4 V, 1:345 / 1:191 / 1:145 |
| Cameras | 2× USB webcam, 640×480 @ 30 fps |
| Host | MacBook Air M1 (MPS backend) |
| Framework | LeRobot 0.6.2 (64b23178) |
| Policy | ACT, 51.6M params, ResNet18 |
| Training | NVIDIA T4 (Kaggle / Lightning AI) |

**Figure 2:** The workspace. The leader arm (black, left) is moved by hand to teleoperate the follower (white, right); the task is performed on the dark mat with the bowl at a fixed, marked position.

**Task.** Pick up a cube from a flat work surface and place it in a bowl. The bowl is fixed; the cube position is randomised across the workspace on every episode.

**Controls.** Every variable except cube position was held fixed within a dataset: camera positions marked on the table and never moved; bowl position marked; constant artificial lighting with blinds closed; fixed arm home position; matte work surface, taped down and outlined.

**Observations.** The policy receives the six joint positions and two camera streams: an overhead view, which shows the object layout, and a front view, which shows gripper height. Their placement defines the two camera setups described next.

## 3 Datasets, setups and models

The project went through three iterations. Each was triggered by a specific failure of the previous one; Table 1 maps datasets to setups and models.

| Dataset | Setup | Episodes | Frames | Protocol | Models |
|---|---|---|---|---|---|
| 1 `so101_cube_demo_0` | 2 | 100 | 44,553 | original | `act_so101_cube_100` |
| 2 `so101_cube_v2` | 2 | 100 | 47,872 | revised | `v2_024`, `v2_100` |
| 3 `so101_cube_v3` | 3 | 51 | 24,404 | revised | `v3_050` |

**Table 1:** Datasets, camera setups and the models trained on them. All datasets are recorded at 30 fps, dataset 1 with 15 s episodes and datasets 2 and 3 with 16 s episodes. `v2_024` uses episodes 0–23 of dataset 2. Models are referred to by short names throughout: `v2_024`, `v2_100` and `v3_050` stand for `act_so101_cube_v2_024`, `act_so101_cube_v2_100` and `act_so101_cube_v3_050`.

### 3.1 Iteration 1: the pilot dataset and a gripper that never opened

The first 100 demonstrations were recorded the way a person naturally performs the task: a quick, loosely controlled reach toward the cube, opening the gripper only on arrival, closing it, and carrying the cube to the bowl. The resulting model (`act_so101_cube_100`) moved toward the cube with a visibly jittery arm and **kept the gripper closed**, or almost closed, throughout. Rather than training longer, the action distribution of the dataset was examined directly.

The gripper is one of six action dimensions, so under MEAN_STD normalisation it contributes one sixth of the loss. Because it opened only in the last instant before the grasp, 61% of all frames had it closed and the median action was 1.2 on a range up to 79 (Figure 3, left). A policy that always predicts "closed" matches the majority of frames and pays almost nothing in loss — while failing at exactly the moment that determines task success.

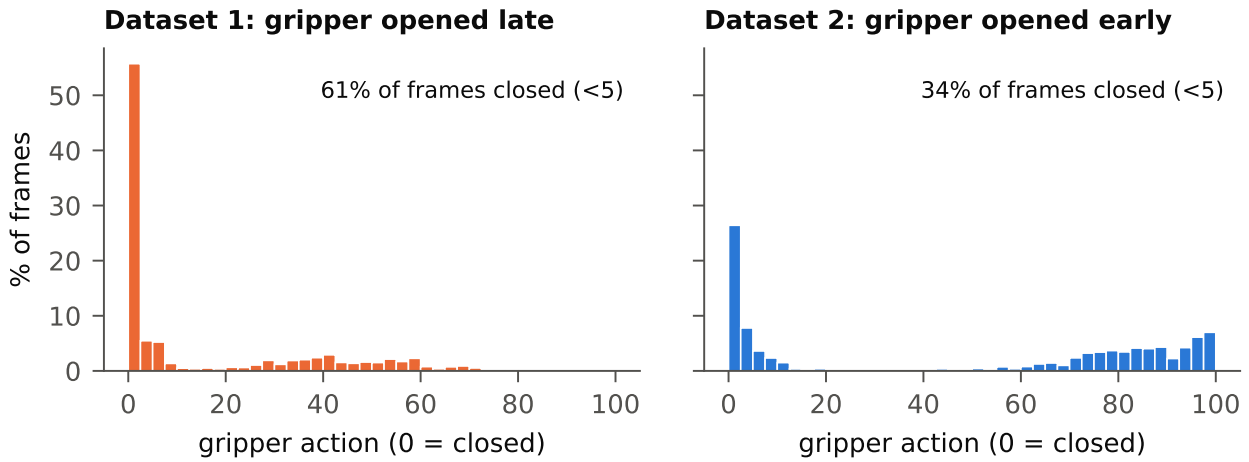



**Figure 3:** Distribution of the gripper action over all frames. With the original protocol (dataset 1, left) the gripper opens only just before the grasp and most frames are "closed". Opening it at the start of the episode (dataset 2, right) moves a large share of frames to the open state.

**Revised protocol.** Dataset 2 is a complete re-recording of all 100 demonstrations, with slower, deliberate demonstrations, often with one hand on the leader's trigger and the other steadying the leader arm, and the gripper opened at the *start* of every episode: home → open gripper fully → slow approach → pause above the cube → decisive close → lift, transport, release above the bowl. Episodes end at release; the return home happens during the unrecorded reset window. The share of "closed" frames fell from 61% to 34%, and every model trained on the revised data opens the gripper correctly. The lesson generalises: in behaviour cloning a failure mode can be a property of the *demonstration distribution* rather than of the model or the training budget, and the dataset is the first place to look.

### 3.2 Setup 2: the main dataset

Setups are numbered to match the dataset and model versions (v2, v3). Datasets 1 and 2 share the same camera arrangement (Figure 4). The overhead camera is fixed under the shelf above the desk and looks straight down on the workspace. The front camera stands on a small tripod at the side of the table, slightly above the work surface, and looks across it toward the arm. The effective workspace — the area in which the cube is placed — is about $40 \times 40$ cm. The object is a pale yellow cube. Dataset 2 is the main dataset of the study.

### 3.3 Setup 3: a perceptual redesign

The models trained on dataset 2 performed poorly (Section 5), and inspection of their observations suggested a perceptual problem (Section 7). To test whether better visibility would help, the scene was redesigned (Figure 5) and dataset 3 recorded with the revised protocol:

1. a saturated green cube instead of the pale yellow one, which had weak contrast against the surface under warm light;

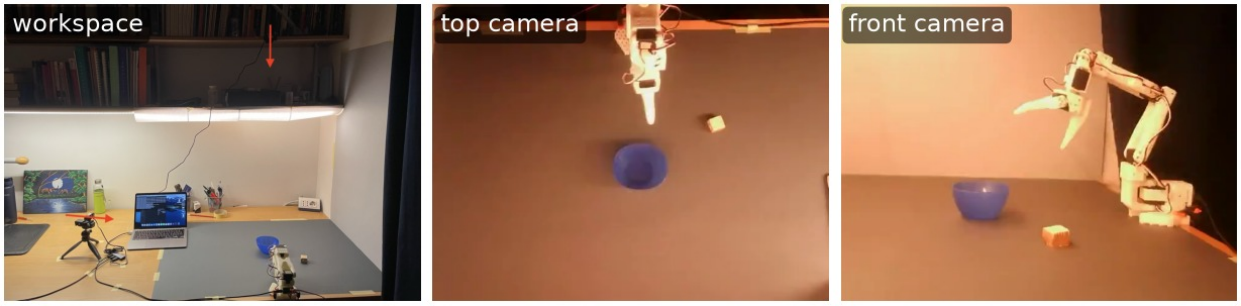



**Figure 4:** Setup 2. Left: the workspace, with arrows marking where the two cameras look. Centre and right: the two observation streams at the resolution the policy receives. The cube occupies roughly 30×30 pixels and is clearly resolved in both views.

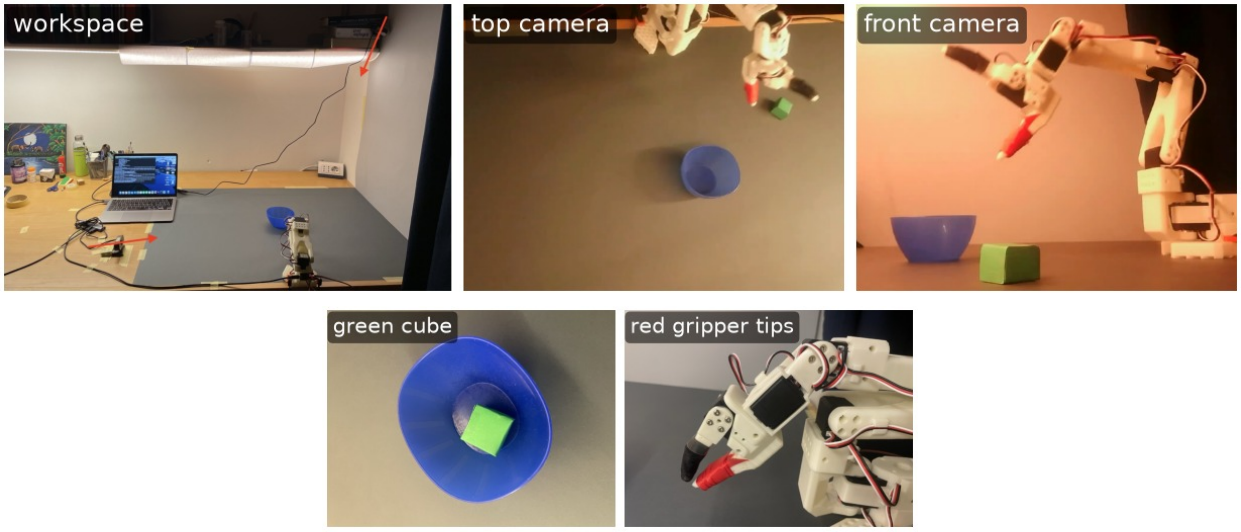



**Figure 5:** Setup 3. Arrows mark where the cameras look: the front camera sits at work-surface level and the overhead camera is tilted by about 30° from the vertical. The green cube and red gripper tips raise contrast in both views.

2. red tape on the gripper tips, a distinct visual reference for the grasp point;
3. the front camera lowered to work-surface level, making gripper height legible;
4. the overhead camera moved to the corner of the shelf and tilted by about 30° from the vertical rather than pointing straight down, recovering perspective;

Space constraints prevented moving the overhead camera to the side opposite the arm's entry, so occlusion during the final descent was reduced but not eliminated.

## 4 Training and evaluation protocol

**Training.** Identical hyperparameters for the three evaluated models: ACT with `chunk_size` = `n_action_steps` = 100, ImageNet-pretrained ResNet18 backbone, AdamW with learning rate $10^{-5}$, batch size 32, 20,000 steps (15,000 for `v3_050`, see Table 2), seed 42. Training cost depends on steps × batch size, not on dataset size: with fewer demonstrations the model simply sees more epochs. The 24 vs. 100 comparison is therefore *data quantity at a fixed optimisation budget* — the question that matters to someone deciding how many demonstrations to record. Because the 24 episodes are a nested subset of the 100, the two models share the same setup and recording sessions; for this to be valid the cube position was randomised episode by episode across the whole workspace.

**Evaluation.** 15 cube positions, marked so that they are visible to the operator but not to the cameras, identical across all models; a fixed rollout duration of 22 s for every model and position; scene conditions identical to data collection. Each rollout is graded on a four-level scale fixed before evaluation:

| | Outcome | Definition | Points |
|---|---|---|---|
| **a** | success | cube ends in the bowl | 3 |
| **b** | partial | cube picked up but not deposited in the bowl | 2 |
| c | near miss | correct motion, but the gripper closes within about 1 cm of the cube | 1 |
| d | failure | anything else | 0 |

The headline metric is the success rate (outcome **a**), reported with Wilson 95% confidence intervals; at $n = 15$ these are wide, and reporting them is the point. The points are a coarse, ordinal summary of *how close* a model gets to the goal. They are not a calibrated measure and are used only to compare the three models qualitatively.

## 5 Results

**Training converges.** All runs reduce $L_1$ loss by an order of magnitude and flatten (Table 2, Figure 6). Optimisation is not the limiting factor.

**More data does not raise success.** Going from 24 to 100 demonstrations moves success from 1/15 to 2/15, with heavily overlapping intervals. The correct reading is **no detectable difference at this sample size**, not that the larger model is better.

**Neither does better visibility.** Setup 3 — higher-contrast object, marked gripper tips, repositioned cameras — also reaches 2/15.

### 5.1 Outcomes by position

Table 3 and Figure 7 show what the success rate hides.

**The dominant outcome is a near miss.** 10, 8 and 12 of the 15 rollouts of the three models end with the gripper closing within about a centimetre of the cube (Figure 1, bottom). The arm lifts, moves to the correct region and opens and closes the gripper in the right order; it fails only in the final alignment.

**A qualitative progression.** Points rise from 13 to 16 to 18. The largest change is not in successes but in complete failures, which drop from 4 (both setup-2 models) to 1 with setup 3: positions 7 and 14, total failures for both setup-2 models, become near misses. With $n = 15$ this difference is not statistically significant (Fisher's exact test, $p = 0.33$), so it is an indication rather than a result — but it points in the same direction as the qualitative observations below.

**Qualitative observations from the evaluation.**

- `v3_050` is generally better at *second attempts*: after a missed grasp it reacts faster and its retry is more accurate.
- Outcomes are not deterministic. A success at a given position does not mean the same position succeeds again; each cell in Table 3 is one sample.
- The policies do not follow the bowl if it is moved: they release at the place where the bowl usually is. Since the bowl was fixed in every demonstration, the transport phase is tied to a coordinate in the workspace rather than to the bowl as an object.

| Model | Demos | Epochs | Final $L_1$ | Success | 95% CI | Points |
|---|---|---|---|---|---|---|
| v2_024 | 24 | 55.7 | 0.050 | 1/15 (6.7%) | 1–30% | 13 |
| v2_100 | 100 | 13.4 | 0.076 | 2/15 (13.3%) | 4–38% | 16 |
| v3_050 | 51 | 19.7† | ≈0.09† | 2/15 (13.3%) | 4–38% | 18 |

**Table 2:** Main comparison. Same hyperparameters, same 15 evaluation positions; points out of 45. †Training of v3_050 stopped at 15,000 of 20,000 steps because of repeated cloud interruptions.

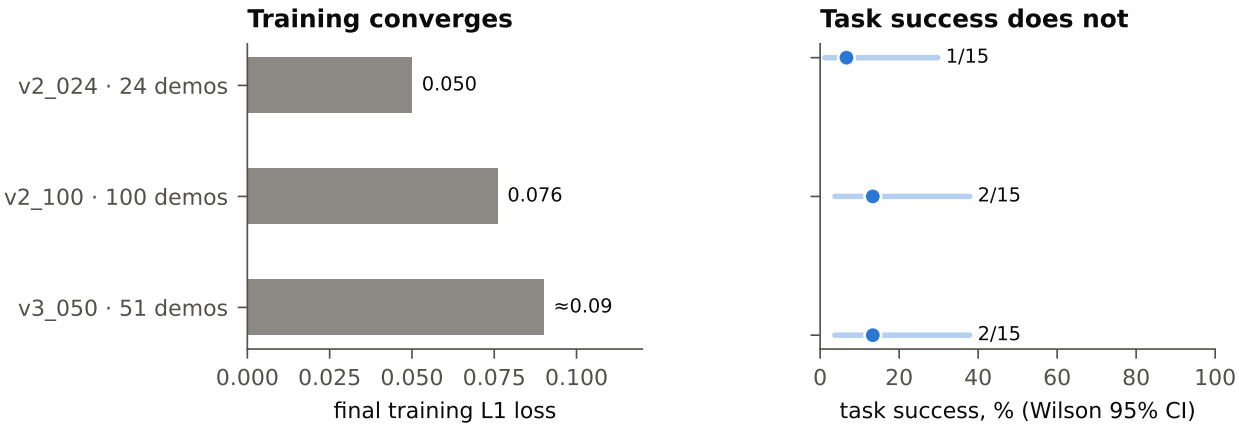



**Figure 6:** Every run reduces the $L_1$ loss by an order of magnitude from ≈0.5–0.6 and flattens (left), yet task success stays low and the confidence intervals overlap almost entirely (right).

| Position | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | Pts |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| v2_024 | c | c | **a** | c | c | c | d | c | c | c | c | d | c | d | d | 13 |
| v2_100 | c | c | **b** | c | **a** | **a** | d | c | d | c | d | c | c | d | c | 16 |
| v3_050 | **a** | c | c | c | c | **a** | c | c | d | c | c | c | c | c | c | 18 |

**a** success **b** picked up, not deposited **c** near miss (≲1 cm) **d** failure (points: a = 3, b = 2, c = 1, d = 0)

**Table 3:** Outcome of every evaluation rollout. Each cell is a single rollout at a fixed, marked cube position.

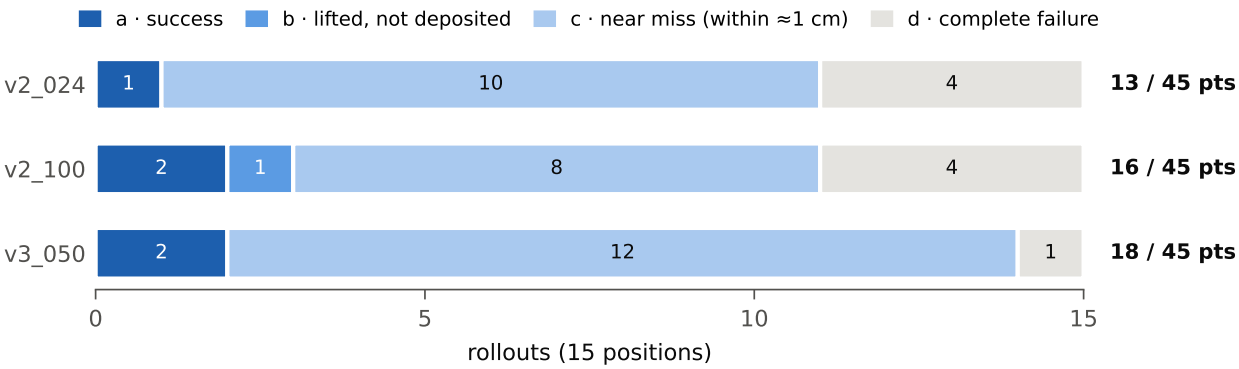



**Figure 7:** Outcome distribution per model. Successes stay at 1–2, but the share of complete failures falls from 4 to 1 with setup 3, and nearly every rollout of v3_050 ends at least in a near miss.

## 6 What the policy did learn

A 7–13% success rate could suggest a policy that has learned little. The rollouts show otherwise. Every evaluated model reproduces the full structure of the task — approach, open, descend, close, lift, transport, release — and aims at the correct region of the workspace for each cube position. The failures are concentrated in the last centimetre.

More tellingly, the policies **react to the outcome of their own grasp**. In additional, informal rollouts with longer durations (30, 40 and 50 s), a missed grasp is often followed by a characteristic sequence: the arm starts moving toward the bowl with an empty gripper, then turns back, returns to the cube and tries again. Figure 8 shows a single, uninterrupted rollout of `v3_050` in which the first grasp misses, the arm leaves, comes back, regrasps and completes the task.




**Figure 8:** One continuous rollout of `v3_050`: the first grasp misses (1), the arm starts toward the bowl with an empty gripper (2), returns to the cube (3), regrasps (4) and deposits it (5). This is an extended-duration rollout outside the evaluation protocol, not one of the 15 counted in Table 3. Video in Section 9.

This behaviour was never demonstrated: every training episode contains exactly one, successful grasp. It means the policy uses its observations to tell whether the cube is in the gripper and where it roughly lies, and re-plans accordingly. The delay before it turns back is consistent with action chunking: with `n_action_steps` = 100 the policy executes about 3.3 s of actions before observing the scene again, so it commits to the transport motion before it can notice the miss.

The picture is therefore not of a policy that ignores vision, but of one that extracts *coarse* information — task phase, grasp success, approximate cube location — and lacks the *fine* information needed to close the last centimetre. That is precisely what the camera geometry predicts.

## 7 Diagnosis: a perceptual bottleneck

**Low loss with poor performance is the key observation.** A low $L_1$ means the policy accurately predicts the *demonstrated* actions. If the exact object position were usable from the observations, conditioning on it would lower the loss further — a trajectory aimed at the actual cube matches each demonstration better than an approximate one. The policy does not get there, which suggests the fine-grained information is not reaching the model in a usable form.

**It is not resolution.** The first hypothesis — that the cube is too small in the input images — was tested by inspecting dataset frames at the resolution the policy receives. The cube is clearly resolved in both views (Figure 4).

**It is geometry.** What is missing is not the object but the *relative* position of gripper and object at the moment that matters (Figure 9). In the overhead view the arm enters from the frame edge and occludes the grasp region as the gripper descends. In the front view gripper height is legible, but depth along the viewing axis is ambiguous. At the instant of closure neither view gives a reliable estimate of the gripper–object alignment error, so the policy has no signal to correct the last centimetre from.

**What setup 3 changed, and what it did not.** Better contrast and camera angles did not raise the success rate, but they may have reduced complete failures and improved retries (Section 5). This is what a geometric bottleneck predicts: improved visibility helps the *coarse* stages — finding the cube, recognising a miss — while the final alignment, which depends on seeing gripper and cube together, is unchanged.

### 7.1 Converging evidence

1. Training loss is low and flat → optimisation is not limiting.
2. 4× more data does not change the success rate → sample count is not limiting.

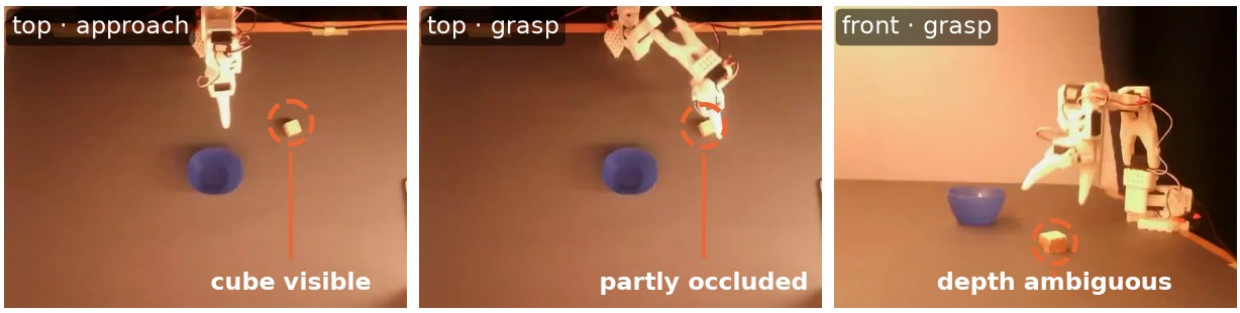



**Figure 9:** The same setup-2 demonstration seen by the policy. During the approach the cube is clearly visible from above (left). At the grasp, the arm covers it in the overhead view (centre), while the front view cannot resolve whether the gripper is in front of or behind it (right).

3. Most failures are near misses, and the policy retries after a miss → task structure and coarse localisation are learned.
4. Better contrast and camera angles do not change the success rate → object *visibility* is not limiting.

What remains is the geometry of what the cameras see at the critical moment.

**The implied fix.** A **wrist-mounted camera** observes gripper and object together, without occlusion, throughout approach and grasp — which is why most manipulation setups include one. It was omitted early in this project for a deformable-object task (where the cloth is occluded by the gripper during grasping) that was later dropped, and the decision was not revisited when the scope changed. Without a 3D printer no adequate mount could be made, and by the time of the diagnosis there was no time to re-record a full dataset. It is the primary next step, and it is motivated by the collected data rather than by general practice.

## 8 Limitations and scope

- **Sample size.** 15 rollouts per model give wide intervals; differences smaller than roughly 25 percentage points in success rate are not detectable.
- **One rollout per position.** Outcomes are not deterministic, so per-position results are single samples, not position-level properties.
- **Ordinal scoring.** The points scale summarises the outcome distribution; it is not a calibrated measure of performance.
- **Two data points, not a curve.** The planned 10/25/50/100 sweep was reduced to 24/100.
- **No out-of-distribution evaluation.** With in-distribution success at 7–13%, OOD evaluation would have produced 0–1 successes out of 15; it was omitted deliberately. The bowl observation in Section 5 is anecdotal.
- **Setup 3 is not a controlled ablation.** Four changes were applied together to maximise the chance of a detectable effect, and `v3_050` training stopped at 15,000 steps.
- **Single operator.** Demonstration quality was not independently assessed.
- **Mechanical repeatability.** The SO-101 uses 3D-printed PLA parts with gearbox backlash; sub-centimetre repeatability was not characterised and may contribute to the near misses.

**Scope management.** The plan was reduced three times, each time deliberately, to guarantee a finished project. Each reduction kept the central question and removed what was redundant or not measurable in the available time.

| | **Original plan** | **Delivered** |
|---|---|---|
| Tasks | rigid cube + deformable cloth | cube only |
| Data levels | 10 / 25 / 50 / 100 | 24 / 100 |
| Rollouts per condition | 25 | 15 |
| OOD factors | position, colour, size, target | none |
| Total rollouts | 320 | 45 |

## 9 Conclusion and future work

The policies trained here succeed on 7–13% of attempts. Four times more data does not raise that rate, and neither does a redesigned, higher-contrast perceptual setup, while training loss converges cleanly in every case. The policies are nonetheless competent at everything except the last centimetre: they reproduce the task structure, aim at the right place, and recover from a missed grasp. Together these locate the limitation: the relative alignment of gripper and object at the moment of grasping is not visible in either camera view when it is needed.

The practical lesson is that camera placement deserves at least as much attention as dataset size, and that for precise manipulation a wrist camera is a structural requirement, not an optional refinement. The methodological lesson is that training loss and task success can diverge sharply in behaviour cloning, that a binary success rate hides much of what a policy has learned, and that the gap between them is where the useful diagnostic information lives.

**Next steps.** In order of expected impact:

1. **Wrist camera** — directly addresses the identified bottleneck.
2. **Varying the bowl position** in the demonstrations, so that placement is conditioned on the bowl rather than on a coordinate.
3. **Reduced workspace** — higher demonstration density for the same episode count.
4. **Non-learning baseline** — OpenCV cube detection plus inverse kinematics, to bound the task and isolate the contribution of learning.
5. **Pre-trained policies** (SmolVLA, $\pi_0$) fine-tuned on the same data, starting from far richer visual representations.
6. **OOD evaluation**, once in-distribution performance is high enough for it to be informative.

## Videos



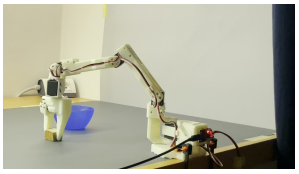
**Project overview** 97 s

From the first successful rollout back to the build, the demonstrations and the evaluated behaviours: success, near miss, and recovery from a missed grasp. Captions mark real-time and accelerated footage.

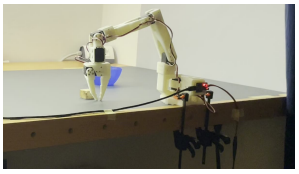
watch the overview
the first success, uncut



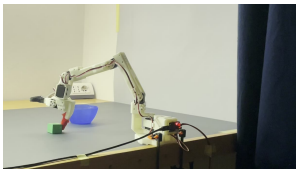
**Teleoperation**
watch video



**Successful rollout**
watch video



**Near miss**
watch video



**Miss, retry, success**
watch video

**Acknowledgements.** Built on the open-source SO-100/SO-101 design by TheRobotStudio and the LeRobot framework by Hugging Face.

## A Reproducing this work

**Published artifacts.**

| Artifact | Identifier |
|---|---|
| Dataset 1, pilot (setup 2) | Fabyo04/so101_cube_demo_0_20260926_155239 |
| Dataset 2, main (setup 2) | Fabyo04/so101_cube_v2 |
| Dataset 3 (setup 3) | Fabyo04/so101_cube_v3 |
| Model, dataset 1 | Fabyo04/act_so101_cube_100 |
| Model, dataset 2, 24 demos | Fabyo04/act_so101_cube_v2_024 |
| Model, dataset 2, 100 demos | Fabyo04/act_so101_cube_v2_100 |
| Model, dataset 3 | Fabyo04/act_so101_cube_v3_050 |
| Code | GitHub repository |

**Environment.** Data collection and inference: LeRobot 0.6.2, commit `64b23178d5348609c266250d3e1f511eba4c33ff`; Python 3.12.13; PyTorch 2.11.0; macOS arm64 (MacBook Air M1). Training: Linux with an NVIDIA T4. Both dataset-2 models (`v2_024` and `v2_100`) were trained with LeRobot 0.6.1, as recorded in their model cards, so the 24 vs. 100 comparison uses a single framework version. LeRobot evolves rapidly, so versions are recorded; results are not expected to reproduce exactly across them.

**Figures.** Every figure in this report is generated by `scripts/make_report_figures.py` from the public datasets, the evaluation sheet (`Report/Evaluation results.xlsx`) and the photos and video stills in `Report/media/`. The stills were extracted once from the original recordings, which are not published, by `scripts/extract_media_stills.py`.

**Data collection.**

```
lerobot-record \
  --robot.type=so101_follower --robot.port=$FOLLOWER_PORT --robot.id=follower_01 \
  --robot.cameras="{ top: {type: opencv, index_or_path: 0, width: 640, height: 480, fps: 30}, \
                front: {type: opencv, index_or_path: 1, width: 640, height: 480, fps: 30} }" \
  --teleop.type=so101_leader --teleop.port=$LEADER_PORT --teleop.id=leader_01 \
  --dataset.repo_id=$HF_USER/so101_cube_v2 --dataset.root=$DATASET_ROOT \
  --dataset.num_episodes=25 --dataset.no_stamp=true \
  --dataset.episode_time_s=16 --dataset.reset_time_s=10 \
  --dataset.streaming_encoding=true --dataset.encoder_threads=2 \
  --dataset.rgb_encoder.vcodec=auto \
  --dataset.single_task="Pick up the cube and place it in the bowl"
```

**Training.**

```
python -m lerobot.scripts.lerobot_train \
  --dataset.repo_id=Fabyo04/so101_cube_v2 \
  --policy.type=act --policy.repo_id=Fabyo04/act_so101_cube_v2_100 \
  --policy.device=cuda \
  --output_dir=outputs/train/act_cube_v2_100 --job_name=act_cube_v2_100 \
  --steps=20000 --batch_size=32 --seed=42 \
  --save_freq=5000 --save_checkpoint_to_hub=true --wandb.enable=false
```

**Autonomous rollout.**

```
lerobot-rollout \
  --robot.type=so101_follower --robot.port=$FOLLOWER_PORT --robot.id=follower_01 \
  --robot.cameras="{ top: {...}, front: {...} }" \
  --policy.path=Fabyo04/act_so101_cube_v2_100 \
  --policy.n_action_steps=100 \
  --device=mps --fps=30 \
  --task="Pick up the cube and place it in the bowl" --duration=22
```

## B Practical notes for building this setup

Things that cost time and are not in the official guides.

**Motor variants are visually identical.** The kit contains four different servos (7.4 V at three gear ratios for the leader, 12 V 1:345 for the follower). Label every motor before assembly. For a spare, the follower's 12 V unit is the right choice: all six follower joints use the same part, and the leader — moved by hand — rarely fails. One vendor datasheet labels the 12 V part "used on leader arm"; this is a copy-paste error. Torque in the product name is the reliable discriminator: 19 kg = 7.4 V = leader, 30 kg = 12 V = follower.

**Set motor IDs one at a time, before assembly.** Servos ship with identical IDs on a shared serial bus. During `lerobot-setup-motors`, disconnect the *DC power*, never the USB — the USB carries the serial port, and unplugging it kills the script with `termios.error: Device not configured`.

**Calibration can silently fail.** A joint reporting a range of 0–4095 has wrapped around the 12-bit encoder boundary, not discovered a 360° joint. Centre every joint in its mechanical travel before the first prompt, then move one joint at a time. Valid ranges stay well clear of both 0 and 4095, and leader/follower amplitudes should match.

**A motor that stops responding is usually in protection.** Power-cycle first. One silent motor suggests a connector; six silent motors suggest upstream power or USB.

**Connectors are the weak point.** Two independent intermittent-contact failures occurred during this project (camera USB cable, leader board USB). Anchor every cable to the table about 10 cm from its connector before starting, not after the second failure.

**Check barrel jack dimensions.** 5.5×2.1 mm and 5.5×2.5 mm have identical outer diameters. A mismatch feels like "needs more force" and will damage the centre pin if forced.

**The frame-rate warning is misleading.** LeRobot's *"Record loop is running slower"* message is emitted only on slow iterations and never reports the mean. With ACT's action chunking, one iteration in a hundred runs full inference and is legitimately slow. Check the frame count in the `Map:` line instead, which should approach `episode_time_s` × `fps`.

**Enable streaming encoding.** Without `--dataset.streaming_encoding=true --dataset.rgb_encoder.vcodec=auto`, video encoding competes with the recording loop. With it, per-episode cycle time dropped from ≈65 s to ≈20 s — the difference between 35 minutes and two hours for 100 demonstrations. Also disable `rerun` visualisation while recording.

**macOS specifics.** Serial ports appear as `/dev/tty.usbmodem<SERIAL>` and are stable across reboots — no udev rules needed. Conversely, `v4l2-ctl` does not exist and the AVFoundation backend generally ignores exposure and white-balance settings; control the lighting instead.

**Test the full pipeline on five episodes first.** Record, train briefly, and run inference end-to-end before committing to a hundred demonstrations.